

## Database Ch. 6 notes

### \* Transforming Data Models into Db Designs \*

#### I. Purpose of a Db Design

Db design: set of db specs that can actually be implemented as a db (DBMS-specific)

#### II. Create a Table for Each Entity

- attr. = ~~co~~
- entity attr = table column
- identifier = table PK

#### Selecting a PK

- facilitates searching and sorting of table rows
- \* Figure 6-1 \* - Transforming Db Model into Db Design
- pg 250

- ideal PK is short, fixed, and numeric (id)
  - use surrogate key if no candidate keys
- disadv of surrogate keys:
  - a) value has no meaning to the user
  - b) when sharing data, if there's 2 db's that have 2 tables w/ same name, they'll have the same ID rows (not likely) (unless merging db's)
    - surrogate keys can have a defined range of values

#### Specifying Alternate Keys

- candidate key: identifier of unique rows w/in a table
- we choose 1 in a table to be the PK
  - the ones not chosen are alternate keys

- representing AK's:



AK 1.1

↳ 1<sup>st</sup> column of the 1<sup>st</sup> alternate key

## Specifying Column Properties

4 types: ~~null~~, ~~not null~~, null status, data type, default value, data constraints

### null status

- column can either be null or not null

### data type

- DBMS-specific
- some are standard across the board
  - Char for fixed-length string
  - Varchar(x) for variable-length string
- companies usually have their own generic data standards

\*pg 254-258 = Common Data Types in MySQL \*

### Default Value

- a value supplied by the DBMS when a new row is made
  - could be string, date/time, timestamp, int, etc.
  - could be the result of a function (ex. calculating price)

Example: table: ITEM

column: ItemPrefix

Default Value: if Category = 'Perishable' then 'P'  
otherwise 'N'

\* Figure 6-8 pg 259

## Data Constraints

- limits on data values (on Ex. only 4 values allowed in this col.)
- Range constraints limit values to a particular interval of values
  - Ex. dates b/w 1/1/1900 to present



- Intrarelation constraint limits column values in comparison w/ other columns in same table
  - Ex. ReviewDate must be at least 3 months after HireDate
- Interrelation constraint limits a column's values in comparison w/ other columns in other tables
  - Ex. DEAD. Name column value cannot equal ALIVE. Name column value
  - Referential Integrity is one type of interrelation constraint

### Verify Normalization

- structures of forms/reports usually reflect how users think about their data
  - will usually end up normalized if modeling from forms/reports
- should verify though

### III. Create Relationships

\* Prerequisite: design a complete, but independent, set of tables

#### Relationships b/w Strong Entities

- generally, relationships are created by placing foreign keys into tables
- 1:1, 1:N, N:M

#### 1:1 Relationships

- ~~option 1~~: place PK of Table1 as FK of Table2, or vice versa
- each value in PK column of T1 can only appear once in T2

- To enforce required uniqueness of the FK value, we can define the FK column as 'unique', either directly or as an alternate key.
  - this is also how we enforce maximum cardinality

### 1:N Relationship

- place PK of T1 on the 1 side (1:N) into the table on the many side (1:N) as a FK
- 1:N  
   ↓     ↓  
   parent, child
  - Don't make the FK unique
- "Place the PK of the parent as the FK of the child"
- \* Figure 6-10 pg 261

### N:M Relationships

- problem: no place in either table to place the FK
  - a 2-way 1:N rel won't work (no duplicates!)
- solution: create an intersection table, holding the PK's of each table only, as FK's
  - \* Figure 6-11 pg 262
- the 2 PK's in the intersection table serve as the composite PK
  - ints. table holds key data only

### Relationships using 1D-Dependent Entities


#### 4 purposes:

- 1.) Repping N:M relationships (intersection table = 1D-Dep)
- 2.) Repping association relationships
- 3.) Storing multivalued attributes
- 4.) Repping archetype/instance relationships



## Association Relationships

- Start:
- intersection table that's ID-Dep on both its parents
  - convert to association table by adding non-identifier attributes



- can connect > 2 tables in intersection table

## Multivalued Attributes

See Figure 15-a and 15-b (pg 267)

## Archetype / Instance Pattern

- if the pattern is given identifiers of their own, it becomes weak but no longer ID-Dependent
  - relationship must then be transformed using rules of a 1:N rel. b/w a strong to weak/non ID-Dep entity
  - means the PK of parent table should be placed in the child table as a FK

## Relationships w/ a Weak, Non-ID Dependent Entity

- relationship b/w 2 strong entities behaves the same
- nonidentifying rel to is characterized by max card
- use surrogate key to replace long identifier situations (parent)
- then replace identifier of ID-Dep entity w/ its own surrogate key

## Relationships in Mixed Entity Designs

Mixed-ID: strong entity to ID-Dependent entity

\* Figure 18 to Figure 19 (pg 269) show great examples of how these relationships work

## Relationships b/w Supertype to Subtype Entities

- 'IS-A' Relationships bc the supertype to subtype represent the same underlying entity
  - \* ~~Ex~~ a MANAGER is an EMPLOYEE
  - \* keys of all subtype tables are same as keys of the supertype
- discriminator attr cannot be rapped in relational designs
  - \* Figure 6-20 (pg 270)

## Recursive Relationships

1:1

- Boxcar data model = Figure 6-21 (pg 271)
- create a FK in child table that contains an identifier of its recursive child
  - \* FK must be unique
- both sides of the relationship are optional (no required children) bc of the first to last boxcar

1:N

- All cases: Place the PK of parent table in child table as a FK
  - if ~~the~~ a data model requires children (i.e. circular), then each parent table has a child to vise versa



## N:M

- trick: decompose N:M into 2 1:N relationships (i.e. create intersection table)

## Representing Ternary & Higher-Order Relationships

- ~~ternaries~~ ternaries could be repped by multiple binary rels (2 tables)
  - issue if a business rule is added that affects the relationship (i.e. particulars)
  - the ability to only allow certain combos in a table can only be enforced by program code
- MUST constraint: requires one entity to be combined w/ another entity
- MUST NOT constraint: combos that aren't allowed in a ternary rel
- MUST COVER constraint: all combos that must appear in a ternary rel
- none of these 3 types of binary constraints can be represented in the relational design

## Relational Representation of Highline Univ. Data Model

- Final data model: \*Figure 6-27 (pg 275)
- Final db design: \*Figure 6-28 (pg 276)
  - look at the design to visualize the tables - it makes sense!
  - note: cannot have 2 columns (PK's & FK's) w/ same name from 2 different tables

## IV. Design for Minimum Cardinality

- complicated bc certain things cant be enforced in db structures
- Rel's can have 4 min. cardinalities: parent / child optional (0-0); parent mandatory & child optional (1-0); 0-1; 1-1

- (0-0) rels don't require enforcement of min. card.
- \* Figure 6-29: actions needed to enforce minimum cardinality
  - (a): When parent is Mandatory
  - (b): When child is mandatory
- "minimum cardinality enforcement action"
- \* actions are for Insert, Update, & Delete functions

### Actions when Parent is Required

- ensure every row has a non-null value for FK column
  - \* restrict action to update or delete parents PK and restrict creating or modifying child FK

### Actions on Parent row when Parent is required

- if update:
  - if any value of parent PK changes, then the child FK becomes an orphan (prohibit modification)
  - \* be consistent in your changes

Cascading Updates: policy of making changes happen from Parent PK to child FK

### - if delete:

- delete PK row = children FK are orphans
  - \* either delete children too or prohibit deletion

Cascading Deletions: deleting the children along w/ the parent

### Actions on Child row when Parent is required

- child must have valid FK value
  - \* could have a default value for the FK
- no restrictions of deleting a child

### Actions when Child is Required

- ensure there's at least 1 child row at all times
  - \* last child cannot leave parent



- enforcing required children is harder than enforcing required parents
  - we must count # children a parent has

### Actions on Parent Row when Child is required

- if child is required, then can't create a new parent w/o creating a relationship to a child
  - either create new child row each time parent is created or change an existing child row's FK
  - else = prohibited
- modification of parent's PK has no restriction if parents use surrogate key
- if child is required & parent is deleted, no action need be taken bc child is required, not the parent

### Action on Child row when Child is required

- no action w/ insertion
- if child is the last of its parent, then update cannot occur; if so, parent would be childless
  - must know # of children at all times
- if last child, cannot delete child row

### Implementing Actions for M-O Relationships

- never create orphans!
- a) - enforce referential integrity so every FK value has a match in a PK table
- b) - make FK column NOT NULL
  - if a to b, then all of \*Figure 6-29 will be enforced

\*Figure 6-31 (pg 280) - Actions to Apply to enforce Minimum Cardinality

## Implementing Actions for O-M Relationships

triggers: modules of code that're invoked by the DBMS when specific events occur

- used for insert, update, delete actions

- use cases:

- parents - create the required parent or child

- steal existing children from parents

- children - on updating/deleting parents, if FK is null, proceed. If FK has a value, and the row is the ~~last~~ last child, the trigger must either:

- Delete the parent

- Find a substitute child

- Disallow (reject) an update/delete

- triggers require code logic

## Implementing Actions for M-M Relationships

- hard to enforce

- Ex: 2 tables; each has an insert trigger; it calls insert on T1, but fails bc trigger requires a row in T2 before an entry in T1 can be made, and vice versa.

- similar issue in deleting

## Designing Special-Case M-M Relationships

- (M-M's) are easier b/w strong & weak entities

- "no one will ever try to insert, update, or delete

- a new child except in the context of the parent

- all actions apply for Action on Parent from Figure 6-29 (a & b)

- insert = always create child

- update/delete = cascade remaining actions



## Documenting Minimum Cardinality Design

### Doc. Required Parents

- 3 design decisions to be made:
  - 1) determining if an update to parent's PK should cascade or be prohibited
  - 2) same as 1, but w/ deletion
  - 3) identifying how a parent row is to be selected on insertion of a child

### Doc. Required Children

- use Figure 6-29 (b) as a boilerplate document (i.e. copy the figure for each rel. w/ a required child & fill in specific actions for C ~~R~~ UD.

## An Additional Complication

- need to specify a design for the minimum card. of every relationship (it'll vary b/w O-M, M-O, & M-M)
  - triggers; some tables will have  $> 1$  for each C ~~R~~ UD function
- \* Figure 6-34 (pg 284) - Summary of design decisions for min. cardinality

## V. View Ridge Gallery Db

- Example design (starts pg. 284)
- the gallery sells a ton of stuff
  - Figure 6-36 (pg 286) has the requirements list
  - Figure 6-37 (pg 287) shows the data model
  - Figure 6-38 (pg 287) shows the db design & Figure 6-39 shows the final design

Figure 6-40 (pg 288) shows all relationships:

		Type	Max	Min
Artist	Work	nonidentifying	1: N	M-O
Work	Trans	" "		M-M
Customer	Trans	identifying		O-O
Customer	Cust-Art	" "		M-O
Artist	Cust-Art	" "	" "	M-O

- Figure 6-41 (pg 289) shows Actions to enforce Minimum Cardinality for required parents
- Figure 6-42 (pg 290) shows " " for Work-To-Trans relationship
- \* - Figure 6-43 (pg 291) shows all column properties

Database Ch. 8 notes: Database Redesign